

Road Guardian AI

YOLOv8n Pothole Detection Model — Accuracy & Performance Report

Generated: 04 September 2026, 00:10

1. Model Overview

Property	Value
Architecture	YOLOv8 Nano (YOLOv8n)
Base Pre-trained Weights	yolov8n.pt (COCO)
Training Strategy	Transfer Learning (pretrained=True)
Training Epochs	30 (converged)
Input Resolution	640 × 640 pixels
Total Parameters	3,011,043
Computational Cost	8.2 GFLOPs
Layers	130
Training Device	CPU
Ultralytics Version	8.4.23
Training Completed	18 March 2026
Confidence Threshold	≥ 60% (live inference)
IoU Threshold (NMS)	0.70

2. Final Validation Metrics (Best Checkpoint — Epoch 29)

The table below shows the metrics recorded on the **held-out validation set** at the best checkpoint epoch. mAP50 is the primary accuracy metric used in the object-detection community and represents mean Average Precision at 50 % Intersection-over-Union (IoU) overlap.

Metric	Value	Interpretation
Precision	79.6 %	Of every predicted pothole, 79.6 % were correct (true positives)
Recall	67.3 %	Of all real potholes, the model successfully detected 67.3 %
mAP @ 50	77.2 %	Primary accuracy: avg precision across all IoU ≥ 0.50 thresholds
mAP @ 50–95	49.8 %	Strict accuracy: avg across IoU 0.50 → 0.95 (industry standard)
Val Box Loss	1.302	Bounding-box regression loss (lower = better localisation)

Val CIs Loss	1.083	Classification loss (lower = fewer misclassifications)
Val DFL Loss	1.301	Distribution Focal Loss for box coordinate refinement
Fitness Score	0.498	Composite weighted score used to select the best checkpoint

3. Training Progression (Epoch-by-Epoch Key Milestones)

Epoch	mAP@50	mAP@50-95	Precision	Recall	Note
1	25.2 %	12.9 %	48.6 %	24.4 %	Baseline — random init from pretrained
5	14.8 %	5.8 %	26.1 %	22.7 %	Warm-up instability
10	60.2 %	32.7 %	74.5 %	49.6 %	Rapid learning phase
15	73.8 %	43.6 %	77.1 %	63.3 %	Model stabilising
20	71.2 %	43.4 %	66.6 %	69.1 %	High recall, precision fluctuating
24	77.2 %	46.8 %	82.7 %	64.2 %	Near-peak performance
25	74.8 %	47.1 %	78.2 %	65.3 %	Slight dip — LR annealing
29	77.2 %	49.8 %	79.6 %	67.3 %	■ BEST checkpoint saved
30	77.5 %	49.8 %	83.0 %	66.1 %	Final epoch — marginal gain

Table 3 — Key epoch milestones extracted from the training checkpoint. Full 30-epoch data available in the training run logs.

4. Training Hyperparameters

Parameter	Value	Parameter	Value
Batch size	16	Optimizer	Auto (AdamW)
Learning rate (lr0)	0.01	LR final (lrf)	0.01
Momentum	0.937	Weight decay	0.0005
Warmup epochs	3	Warmup momentum	0.8
Mosaic augment	1.0 (on)	Flip LR	0.5
Scale	0.5	Translate	0.1
HSV Hue	0.015	HSV Saturation	0.7
HSV Value	0.4	Erasing	0.4
Box loss weight	7.5	CIs loss weight	0.5
DFL loss weight	1.5	Close mosaic epoch	10
AMP (mixed prec.)	True	Seed	0

5. Live Inference Configuration

Setting	Value	Effect
Model format	ONNX (OpenCV DNN)	Lightweight, no PyTorch needed at runtime
Fallback model	.pt via Ultralytics	Used if ONNX loading fails
Confidence threshold	≥ 60 %	Only high-confidence detections accepted
IoU NMS threshold	0.70	Non-max suppression aggressiveness
Accepted classes	pothole, damage, crack	Other classes are filtered out
Min bbox size	> 0 × 0 px	Degenerate boxes discarded
Image validation	≥ 100 × 100 px	Blurry/tiny uploads rejected
Deduplication	pHash + GPS proximity	Prevents duplicate hazard records

6. Recommendations for Accuracy Improvement

#	Recommendation	Expected Gain
1	Train for 50–100 epochs (model still improving at epoch 30)	+2–4 % mAP
2	Switch to YOLOv8s or YOLOv8m (larger backbone)	+5–10 % mAP
3	Train on GPU with larger batch size (32–64)	Faster + better augmentation
4	Add night/rain/wet-road images to dataset	Better real-world robustness
5	Use CutMix / MixUp augmentation (currently off)	+1–2 % mAP
6	Lower confidence threshold to 0.45 to increase recall	+8 % recall